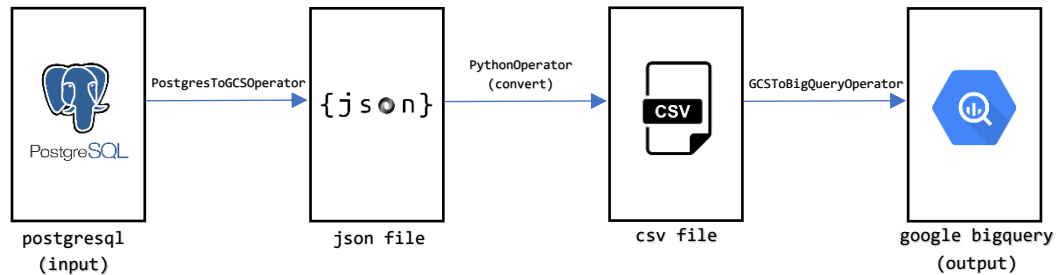


Interview Challenge (Data Engineer) – Document

Apache Airflow Architecture



Dags

- `postgresql_to_gcs_user_log.py` – Dag file that query 'user_log' data from postgresql and upload to GCS
 - `postgresql_to_gcs_users.py` – Dag file that query 'users' data from postgresql and upload to GCS
 - `user_log_to_bigquery.py` – Dag file that transform 'user_log' postgresql data to csv file and load to Google BQ
 - `users_to_bigquery.py` – Dag file that transform 'users' postgresql data to csv file and load to Google BQ
-

Google Cloud Platform Service Usage

- Google **Compute Engine**
 - Google **Storage**
 - Google **BigQuery**
 - Google **Dataproc**
-

airflow-vm - GCE instance

IP: 35.213.169.173

Machine Type: e2-highcpu-4 (4 vCPUs, 4 GB memory)

Description: Apache Airflow Single Node

Airflow URL: <http://35.213.169.173:8080/home>

Airflow Account:

- Admin
 - Username: jon
 - Password: password
- Viewer
 - Username: bluepi
 - Password: password

Airflow PATH: /srv/airflow

Airflow DAG(s) Directory: /srv/airflow/dags

Services:

1. airflow-webserver.service - autostart
2. airflow-scheduler.service - autostart

Usage Guide:

1. `sudo su airflow`
2. `source /srv/airflow/bin/activate`
3. `sudo systemctl status airflow-webserver.service`
4. `sudo systemctl status airflow-scheduler.service`

jupyter-bigquery-m - GCE instance

IP: 35.213.131.149

Machine Type: n1-standard-4 (4 vCPUs, 15 GB memory)

Description: Jupyter Notebook master node for Google Dataproc Cluster

jupyter-bigquery-w-0 - GCE instance

IP: 35.213.135.153

Machine Type: n1-standard-4 (4 vCPUs, 15 GB memory)

Description: Jupyter Notebook worker node 1 for Google Dataproc Cluster

jupyter-bigquery-w-1 - GCE instance

IP: 35.213.143.219

Machine Type: n1-standard-4 (4 vCPUs, 15 GB memory)

Description: Jupyter Notebook worker node 2 for Google Dataproc Cluster

airflow-postgres - GCS instance

Description: Contain JSON and CSV file for Airflow Pipeline

bigquery_bluepi_output - GCS instance

Description: Contain files that use in Google Dataproc Cluster

dataproc-temp* - GCS instance

Description: Contain files when Google Dataproc Cluster get processing

sirapob-bluepi-de-exam:airflow_gcs_to_bigquery - Google BigQuery

Description: Google BigQuery Dataset

sirapob-bluepi-de-exam:airflow_gcs_to_bigquery.user_log_to_bigquery - Google BigQuery

Description: Google BigQuery 'user_log' Table

sirapob-bluepi-de-
exam:airflow_gcs_to_bigquery.users_to_bigquery - Google
BigQuery

Description: Google BigQuery 'users' Table

jupyter-bigquery - Google Dataproc

Jupyter Notebook URL: [Jupyter Notebook](#)

Type: Dataproc Cluster

Cluster Detail:

1. jupyter-bigquery-m - Master
 2. jupyter-bigquery-w-0 - Worker
 3. jupyter-bigquery-w-1 - Worker
-

Jupyter Notebook

Output Monitor Files

- pySpark_user_log_monitor.ipynb
- pySpark_users_monitor.ipynb

pySpark_user_log_monitor.ipynb – Jupyter Notebook IPYNB

```
In [2]: from pyspark.sql import SparkSession
spark = SparkSession.builder \
    .appName('pySpark_user_log_monitoring') \
    .config('spark.jars', 'gs://spark-lib/bigquery/spark-bigquery-latest.jar') \
    .getOrCreate()

In [3]: spark.conf.set("spark.sql.repl.eagerEval.enabled", True)

In [4]: table = "sfrapob-bluepi-de-exam:airflow_gcs_to_bigquery.user_log_to_bigquery"

user_log_data = spark.read \
    .format("bigquery") \
    .option("table", table) \
    .option("dateFormat", "yyyy-MM-dd HH:mm:ss") \
    .load()

user_log_data.printSchema()

root
 |-- action: string (nullable = false)
 |-- created_at: timestamp (nullable = false)
 |-- id: string (nullable = false)
 |-- success: boolean (nullable = false)
 |-- updated_at: timestamp (nullable = false)
 |-- user_id: string (nullable = false)

In [5]: user_log_data_table = user_log_data \
    .select("id", "user_id", "action", "success", "created_at", "updated_at")

user_log_data_table.toPandas()

Out[5]:
```

	id	user_id	action	success	created_at	updated_at
0	4b796e06-3178-4133-ad30-dc505bfc13f5	55514cf0-3026-404f-8ea3-f41b00bdf6b5	login	False	2020-02-17 01:46:05.934519	2020-02-17 01:46:05.934519
1	8745cacb-f8aa-4294-b824-2d3a5c50f171	55514cf0-3026-404f-8ea3-f41b00bdf6b5	login	False	2020-02-17 01:46:06.934519	2020-02-17 01:46:06.934519
2	c12bcaa2-563c-416e-91db-36d846b0feae	55514cf0-3026-404f-8ea3-f41b00bdf6b5	change password	False	2021-03-22 02:10:15.010466	2021-03-22 02:10:15.010466
3	6cfd027d-b50b-4855-b279-91f99ce4476d	55514cf0-3026-404f-8ea3-f41b00bdf6b5	change password	True	2020-02-17 01:50:07.934519	2020-02-17 01:50:07.934519
4	c0c97762-694b-45fd-a41e-7304313eab82	55514cf0-3026-404f-8ea3-f41b00bdf6b5	login	True	2020-02-17 01:56:07.934519	2020-02-17 01:56:07.934519
5	59131a90-62c9-40d8-899e-31c3dee8ad7e	d0e73a35-ff6a-4f64-89b4-ed2b813782a3	login	True	2020-02-17 01:48:08.934519	2020-02-17 01:48:08.934519
6	5bccf37b-5d03-4c54-a76e-5eb54a3290a5	d0e73a35-ff6a-4f64-89b4-ed2b813782a3	logout	True	2020-02-17 01:52:07.934519	2020-02-17 01:52:07.934519

```
In [35]: spark.conf.set("spark.sql.execution.arrow.enabled", "true")

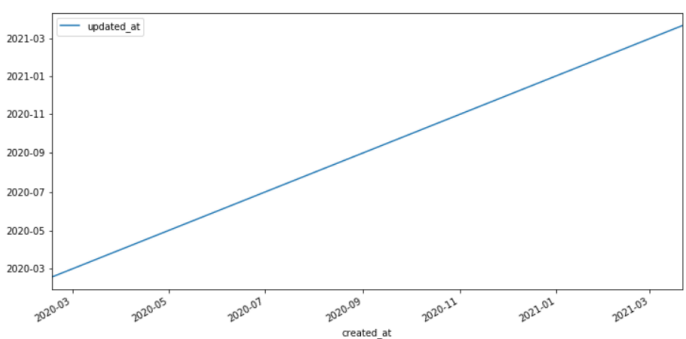
user_log = user_log_data_table.toPandas()
user_log.set_index('created_at', inplace=True)
user_log.head()

Out[35]:
```

	id	user_id	action	success	updated_at
created_at					
2020-02-17 01:46:05.934519	4b796e06-3178-4133-ad30-dc505bfc13f5	55514cf0-3026-404f-8ea3-f41b00bdf6b5	login	False	2020-02-17 01:46:05.934519
2020-02-17 01:46:06.934519	8745cacb-f8aa-4294-b824-2d3a5c50f171	55514cf0-3026-404f-8ea3-f41b00bdf6b5	login	False	2020-02-17 01:46:06.934519
2021-03-22 02:10:15.010466	c12bcaa2-563c-416e-91db-36d846b0feae	55514cf0-3026-404f-8ea3-f41b00bdf6b5	change password	False	2021-03-22 02:10:15.010466
2020-02-17 01:50:07.934519	6cfd027d-b50b-4855-b279-91f99ce4476d	55514cf0-3026-404f-8ea3-f41b00bdf6b5	change password	True	2020-02-17 01:50:07.934519
2020-02-17 01:56:07.934519	c0c97762-694b-45fd-a41e-7304313eab82	55514cf0-3026-404f-8ea3-f41b00bdf6b5	login	True	2020-02-17 01:56:07.934519

```
In [36]: user_log.plot(kind='line', figsize=(12,6))

Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x7f666e468d30>
```



pySpark_users_monitor.ipynb - Jupyter Notebook IPYNB

```
In [1]: from pyspark.sql import SparkSession
spark = SparkSession.builder \
    .appName('pySpark_users_monitoring') \
    .config('spark.jars', 'gs://spark-lib/bigquery/spark-bigquery-latest.jar') \
    .getOrCreate()
```

```
In [2]: spark.conf.set("spark.sql.repl.eagerEval.enabled", True)
```

```
In [3]: table = "sirapob-bluepi-de-exam:airflow_gcs_to_bigquery.users_to_bigquery"
```

```
users_data = spark.read \
    .format("bigquery") \
    .option("table", table) \
    .option("timestampFormat", "yyyy-MM-dd HH:mm:ss") \
    .load()
```

```
users_data.printSchema()
```

```
root
|-- created_at: timestamp (nullable = false)
|-- first_name: string (nullable = false)
|-- id: string (nullable = false)
|-- last_name: string (nullable = false)
|-- updated_at: timestamp (nullable = false)
```

```
In [4]: users_data_table = users_data \
    .select("id", "first_name", "last_name", "created_at", "updated_at")
users_data_table.toPandas()
```

```
Out[4]:
```

	id	first_name	last_name	created_at	updated_at
0	55514cf0-3026-404f-8ea3-f41b00bdf6b5	John	Henry	2020-02-17 01:33:57.796067	2020-02-17 01:33:57.796067
1	d0e73a35-ff6a-4f64-89b4-ed2b813782a3	สมบุญดี	รุ่งแก้ว	2020-02-17 01:33:57.796067	2020-02-17 01:33:57.796067
